

✔ Congratulations! You passed!

Grade received 100% To pass 80% or higher

Go to next item

1. To get the most value out of your machine learning models, you need to deploy them into production to start making practical decisions quickly. Which of the following are the most important factors to consider when choosing a deployment option?

1 / 1 point

✔ Network connectivity



Correct

That's right! This is important because some deployment options require access to a network or internet.

❑ Model Accuracy

✔ Latency



Correct

Correct! The response time of the deployed model is very essential to the choice of deployment.

✔ Cost



Correct

Correct! Cost is definitely something to consider when choosing a deployment option depending on your use case (ie. real-time, batch, edge, etc)

2. As a machine learning engineer on the marketing team, you want to serve advertisements to a customer based on their current session activity. For a better user experience, you want to serve the ads as quickly as possible. What is the most ideal cloud deployment option for a model of this type?

1 / 1 point

❑ Edge deployment

● Real-time deployment

❑ Batch inference deployment

❑ All of the above



Correct

Correct! A real-time deployment option is ideal for this use case.

3. You would like to deploy a model that detects pneumonia from x-ray images of patients in a rural area. Internet connection is a major challenge here and can slow down the process.

1 / 1 point

True/False: Edge deployment and batch inference deployment are both great options for deploying this model.

❑ True

● False



Correct

Correct! Edge deployment works well in cases where there is limited or no network connectivity while the batch inference pipeline requires some form of network connectivity to work properly.

4. When selecting a strategy to deploy a new or updated model, you want to minimize risk, and reduce downtime. Additionally, we need to measure and compare performance relative to the original model. Which of the following statements correctly describes the canary deployment strategy.

1 / 1 point

❑ All users are shifted from the older version to the new, replacement version all at once.

● A small percentage of users are initially exposed to the new version of the model while the majority of users continue using the original version of the model.

❑ A new version of the model is deployed into production alongside a live model, and traffic is sent to both models simultaneously.

❑ None of the above



Correct

Correct! This is called the “canary deployment” strategy, when sets of users are incrementally exposed to the new model version.

5. In both A/B testing and canary deployment strategies, you split your traffic to compare different model versions. What is the difference between these 2 strategies?

1 / 1 point

● A/B testing targets larger groups and typically runs for longer periods of time.

❑ A/B testing deployment is more dynamic and uses reinforcement learning to dynamically shift traffic to the winning model versions.

❑ Canary deployment is more dynamic and uses reinforcement learning to dynamically shift traffic to the winning model versions.

❑ Canary deployment targets larger groups and typically runs for longer periods of time.



Correct

Correct! A/B testing needs to run over longer periods of time to gather performance data that are statistically significant.

6. Amazon SageMaker provides a feature that allows you to host a sequential chain of models behind a single endpoint. What is this feature called?

1 / 1 point

❑ SageMaker CloudWatch

❑ Multi-model Endpoints

❑ Autoscale

● SageMaker Inference Pipeline



Correct

Correct! This allows you to perform data transformations before and after the model prediction in a sequence of steps behind a single endpoint.

7. Amazon SageMaker Model Monitor uses the open source AWS Deequ library to monitor when inference input data drifts away from the baseline training input data. What type of data drift is Model Monitor detecting in this case?

1 / 1 point

● Data quality

❑ Statistical bias drift

❑ Model quality

❑ Feature attribution drift



Correct

Correct! Data quality is used to monitor signals that the current input data coming in for inference is statistically different from the feature data used to train the model.